

General Regulations.

- Please hand in your solutions in groups of three people. A mix of attendees from different tutorials is fine.
- Your solutions to theoretical exercises can be either handwritten notes (scanned), typeset using L^AT_EX, or directly in the jupyter notebook using Markdown.
- For the practical exercises, the data and a skeleton for your jupyter notebook are available at <https://github.com/heidelberg-hepml/mlph2023-Exercises>. Always provide the (commented) python code as well as the output, and don't forget to explain/interpret the latter, we do not give points for code that does not run. Please hand in both the notebook (.ipynb) and an exported pdf.
- Submit all your files in the Übungsgruppenverwaltung, only once for your group of three.

1 KL-divergence

Compute the KL-divergence of two Gaussians

$$D_{\text{KL}}[\mathcal{N}_{\mu_1, \sigma_1}(x), \mathcal{N}_{\mu_2, \sigma_2}(x)], \quad \mathcal{N}_{\mu, \sigma}(x) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right). \quad (1)$$

What happens in the case $\mu_1 = \mu_2, \sigma_1 = \sigma_2$? (2 pts)

2 Random numbers

- (a) Given two probability density functions f and g and a sample $r \sim f$, one can generate another sample $x \sim g$ using the relation

$$x = G^{-1}(F(r)), \quad (2)$$

where F, G are the cumulative density functions of f, g . Derive this relation. (1 pts)

- (b) Use Eq. 2 to generate random numbers from the distributions g_i using $r \sim \text{Uniform}[0, 1]$. If it is not possible, explain why.

- $g_1(x) = 3x^2, x \in [0, 1]$
- $g_2(x) = e^{-x}, x > 0$
- $g_3(x) = \mathcal{N}_{0,1}(x)$
- $g_4(x) = \frac{2}{3\zeta_3} \frac{x^2}{e^x + 1}, x > 0$

Hint: `scipy.special.erfinv`, `scipy.special.zeta` (2 pts)

- (c) Repeat exercise (b) using rejection sampling. Restrict the region to $x < |5|$ whenever it extends to infinity. Compare the speed to your implementation from (b). Discuss pros and cons of this method, compared to the method from (a) and (b). (2 pts)

- (d) These methods for generating gaussian random numbers are too slow. Can you do better? Hint: Box-Muller

(1 pts)

3 Neyman-Pearson Lemma

The Neyman-Pearson Lemma is a result from hypothesis tests that will show up a lot when constructing losses for neural networks. The goal of this exercise is to prove the Neyman-Pearson Lemma.

We compare two hypotheses, the null hypothesis H_0 and an alternative hypothesis H_1 . They are quantified through statistical models $p(x|H_0), p(x|H_1)$. We use the test statistic $t(x)$ to accept/reject points if $t(x) < t_0/t(x) > t_0$ with some cutoff t_0 . The rejection region ω is the set of all rejected points $\omega = \{x|t(x) > t_0\}$. This allows us to define two basic properties of a hypothesis test, the size $p(\omega|H_0)$ and the power $p(\omega|H_1)$.

The Neyman-Pearson Lemma states that the likelihood ratio $t_{\text{LR}}(x)$ is the optimal test statistic. This means that $t(x)$ has maximal power $p(\omega|H_1)$ for a test with given size $p(\omega|H_0)$. The likelihood ratio is defined as

$$t_{\text{LR}}(x) = \frac{p(x|H_1)}{p(x|H_0)}. \quad (3)$$

Prove the Neyman-Pearson Lemma by showing that any modification of the rejection region does not increase the power of the test.

Hint: Modify the rejection region by adding A while removing B , where the regions A and B have to be chosen such that the size of the test remains constant. Then look at the change in power $p(A|H_1) - p(B|H_1)$ that this modification gives. (2 pts)